

Round B: internal pattern descriptions from AI coding episodes

About 30-40 minutes. 28 items. No trick questions and no penalty for guessing.

What you're looking at

For 28 separate interactions between an AI coding agent and a task, a program looked inside the model at one or more moments and tried to write, in English, what pattern of internal activity it found. It did not see the original conversation -- only a numerical pattern, which it turned into a description.

Each item below shows the description(s) generated for one interaction. Most items have three descriptions (taken at three different moments in the same interaction); a few have only one.

A note on the writing

These descriptions were generated automatically and are sometimes awkward, repetitive, or offer several different guesses for the same missing word in the same sentence. That is a property of how they were produced, not something to correct -- read past it and answer based on what's actually there.

Your task: two questions per item

Question 1. Every interaction was run under one of two setups. We are calling them Group 1 and Group 2 and are not telling you what differs between them. Read the description(s) for an item and guess: Group 1 or Group 2? There is no way to know for certain -- go with whatever the writing suggests to you, or guess if nothing suggests anything. Also rate how confident you are, from 1 (pure guess) to 5 (very sure).

Question 2. Does the description talk about the model giving up, falling back to some kind of excuse or placeholder, or admitting that it could not finish something? Answer yes or no.

Please do not look anything up or paste this into an AI assistant. We want your own read of the text, unaided.

When you're done

Fill in `round_b_answers.csv` and send it back. The last rows are for your name, the date, and roughly how long it took.